

● JANUARY 2026 SERIES

# FROM GO BUILD TO GO RUN

GOLANG 2026 - NIV RAVE

## #37

# MUTEX VS. RWmutex

PROTECTING STATE WITHOUT KILLING THROUGHPUT





# The Mutual Exclusion (Mutex)

## The *sync.Mutex* Contract

A standard *sync.Mutex* is binary. It is either locked or unlocked. If one goroutine holds the lock, everyone else (both readers and writers) must wait in line.

```
● ● ●

type Counter struct {
    mu    sync.Mutex
    value int
}

func (c *Counter) Inc() {
    c.mu.Lock()
    defer c.mu.Unlock() // Always defer immediately
    c.value++
}
```

Standard Mutexes are incredibly fast and should be your default. They have very low overhead. Use them when your "critical section" (the code between Lock and Unlock) is extremely short.





# The Reader-Writer Mutex

## Scaling the Reads

In many systems, like a configuration map or a session cache, 99% of the operations are reads and only 1% are writes. A standard Mutex is wasteful here because readers don't actually interfere with each other.

### The Solution:

`sync.RWMutex` allows unlimited concurrent readers OR one single writer.





# RLock vs. Lock

## Implementing the RWMutex

Use *RLock()* when you are only reading. It allows other readers to proceed. Use *Lock()* when you are modifying. It blocks everyone.

```
● ● ●

type SessionCache struct {
    mu    sync.RWMutex
    data map[string]string
}

func (s *SessionCache) Get(key string) string {
    s.mu.RLock()
    defer s.mu.RUnlock()
    return s.data[key]
}

func (s *SessionCache) Set(key, val string) {
    s.mu.Lock()
    defer s.mu.Unlock()
    s.data[key] = val
}
```





# The RWMutex Performance Trap

Is *RWMutex* always faster?

**Counter-intuitive Fact:** A *RWMutex* is more complex than a standard *Mutex*. It has to maintain a reader count and handle "writer starvation" prevention logic.

If your read operation is tiny (e.g., just returning a struct field), a standard *Mutex* is often faster because its internal overhead is lower. Only switch to *RWMutex* if your profiling shows significant contention from readers.





# No Reentrancy (The Deadlock)

Recursive Locking is Fatal

Go's Mutexes are not reentrant (unlike Java's *ReentrantLock*). If a goroutine tries to *Lock()* a Mutex it already holds, it will deadlock itself forever.

```
● ● ●

type Counter struct {
    mu    sync.Mutex
    value int
}

func (c *Counter) Increment() {
    c.mu.Lock()
    defer c.mu.Unlock()
    c.value++
}

func (c *Counter) Add(n int) {
    c.mu.Lock()
    defer c.mu.Unlock()
    c.Increment() // DEADLOCK! Increment tries to Lock mu again.
}
```

The Fix: Move the logic into an unexported helper (e.g., *increment()*) that doesn't handle locking, and have both public methods call it.





# Copying Locks is Forbidden

## Passing the State, not the Lock

A `sync.Mutex` contains internal state (a waiter count). If you copy a struct that contains a Mutex, you copy that state.

Always pass structs with locks by pointer (`func (c *Counter)...`).



```
type Counter struct {
    mu    sync.Mutex
    value int
}

// WRONG: Passed by value.
// Every call creates a NEW copy of the mutex. No actual protection!
func (c Counter) BadValue() int {
    c.mu.Lock()
    defer c.mu.Unlock()
    return c.value
}

// RIGHT: Pointer receiver.
// All calls share the same mutex instance.
func (c *Counter) GoodPointer() int {
    c.mu.Lock()
    defer c.mu.Unlock()
    return c.value
}
```





# Lock Contention & Granularity

## The "Big Lock" Anti-Pattern

Locking your entire application state with one global Mutex is a performance killer.

### **The Senior Move:**

Fine-grained Locking. Instead of one lock for a map of 1,000,000 items, use a "Sharded Map" where 32 separate Mutexes protect different segments of the data. This drastically reduces the probability of two goroutines colliding.





# Summary:

- *Mutex*: Fast, simple, default choice.
- *RWMutex*: Best for read-heavy, slow-read operations.
- No Reentrancy: Never lock twice in the same goroutine.
- Pointers only: Never copy a Mutex.

**Question: Have you ever seen an RWMutex actually slow down an app compared to a regular Mutex? Why do you think that happened? ↗**

