

● JANUARY 2026 SERIES

# FROM GO BUILD TO GO RUN

GOLANG 2026 - NIV RAVE

## #41

# TIMERS AND TICKERS: THE PULSE OF GO

MASTERING TIME-BASED EVENTS IN CONCURRENT SYSTEMS





# Beyond time.Sleep

## Why Sleep is Not Enough

*time.Sleep* blocks the entire goroutine, making it impossible to respond to other signals (like a cancellation) while waiting.

### The Concept:

To build responsive systems, we need non-blocking ways to wait. By using Channels with Timers, we can "listen" for a timeout while simultaneously listening for work or shutdown signals.





# Precision Timing with `time.Timer`

## The One-Shot Timer

A `time.Timer` sends a signal on its channel `C` exactly once after a specified duration. Unlike a simple sleep, a Timer can be stopped or reset before it fires.



```
// Initialize a timer for 2 seconds
timer := time.NewTimer(2 * time.Second)

// This blocks until the timer's channel receives the signal
// It's a more flexible alternative to time.Sleep
expirationTime := <-timer.C
fmt.Println("Timer expired at:", expirationTime)
```





# Timer in a Select

## Non-blocking Waits

The true power of a Timer is realized when paired with a `select` statement. This allows your goroutine to remain responsive to multiple sources.

```
● ● ●  
timer := time.NewTimer(5 * time.Second)  
defer timer.Stop()  
  
select {  
case <-timer.C:  
    fmt.Println("Timer fired!")  
case <-ctx.Done():  
    fmt.Println("Operation cancelled; stopping timer.")  
    if !timer.Stop() {  
        <-timer.C // Drain the channel if it already fired  
    }  
}
```





# The Heartbeat Pattern

Recurring Events with *time.Ticker*

While a Timer fires once, a *time.Ticker* fires repeatedly at a set interval. This is the foundation for background tasks, health checks, and periodic cache flushes.

```
● ● ●

// Create a ticker that ticks every 500 milliseconds
ticker := time.NewTicker(500 * time.Millisecond)

// Use a loop to receive values from the ticker channel
for i := 0; i < 3; i++ {
    t := <-ticker.C
    fmt.Println("Tick at", t)
}

// Crucial: stop the ticker to release associated resources
ticker.Stop()
```

Always remember to *Stop()* your tickers. A leaked ticker will continue to attempt to send values to its channel, wasting CPU cycles and memory even if the loop using it has exited.





# The Ticker Loop

## Orchestrating Background Work

Use a ticker inside a for-select loop to create a robust background worker that responds to system shutdowns.



```
ticker := time.NewTicker(1 * time.Minute)
defer ticker.Stop()

for {
    select {
    case <-ctx.Done():
        return // Graceful exit
    case t := <-ticker.C:
        // Triggered every minute
        doPeriodicCleanup(t)
    }
}
```





# The Callback Alternative

Efficiency with *time.AfterFunc*

Sometimes you don't want to manage a channel at all. *time.AfterFunc* runs a specific function in its own goroutine after a delay.

## The Use Case:

This is highly efficient for "fire-and-forget" delayed logic, such as logging a warning if a process takes too long without blocking the main execution path.

```
● ● ●  
  
// Log a warning if the operation isn't done in 2 seconds  
timer := time.AfterFunc(2*time.Second, func() {  
    log.Println("Warning: Operation is exceeding threshold")  
})  
defer timer.Stop()  
  
performHeavyWork()
```





# The Performance Trap

*time.After() vs. time.NewTimer()*

*time.After(d)* is a convenient shortcut, but it has a hidden cost in high-frequency loops.

## The Pattern:

*time.After* creates a new *Timer* every time it's called, but you cannot stop it manually. In a tight loop, this can lead to thousands of "ghost" timers sitting in memory until they expire.

## Best Practice:

For long-running loops or performance-critical paths, always use *time.NewTimer* and manually *Stop()* or *Reset()* it.





## Summary:

- Use Timers for one-off delays that need to be cancellable.
- Use Ticklers for recurring "heartbeat" tasks.
- Use AfterFunc for efficient, channel-free callbacks.
- Avoid time.After in long-running loops to prevent memory bloat.

**Next, we scale our concurrency: Worker Pools and the Fan-in / Fan-out pattern.**

