

● JANUARY 2026 SERIES

# FROM GO BUILD TO GO RUN

GOLANG 2026 - NIV RAVE

#55

## MASTERING THE DATABASE: POSTGRES IN GO

HIGH-PERFORMANCE RELATIONAL PERSISTENCE WITH PGX AND  
SQLX





# Why We Avoid the "Magic" ORM

## SQL is Your Superpower

In other ecosystems, the ORM is king. In Go, we value transparency. Heavy ORMs often generate inefficient queries, hide performance bottlenecks, and make complex joins a nightmare to debug.

### The Philosophy:

By writing SQL, you stay in control of your indexes and query plans. We use "SQL Builders" or "Scanners" rather than full ORMs. This keeps your data layer thin and your database happy, especially when dealing with high-throughput production loads.



# The Postgres Specialist - pgx

## Performance Beyond database/sql

Go's standard library *database/sql* is a great general-purpose interface, but it has limitations when you want to squeeze every drop of performance out of Postgres.

Enter *pgx*.

### The Advantage:

*pgx* is a pure-Go driver that supports Postgres-specific features like *Binary Format*, *COPY* for fast bulk loads, and native support for *JSONB*. It is significantly faster and more feature-complete than the standard driver for Postgres-heavy applications.





# Mapping with sqlx

## Reducing the Scanner Boilerplate

The most tedious part of raw SQL in Go is scanning rows: `err := rows.Scan(&u.ID, &u.Name, ...)`. If you have a table with 30 columns, this is a recipe for bugs. `sqlx` solves this by providing "Struct Tag" mapping.

### The Concept:

`sqlx` is a lightweight wrapper around the standard library. It allows you to "Get" a single row or "Select" a slice of rows directly into your structs, matching database columns to struct fields automatically via tags.

```
● ● ●  
  
type User struct {  
    ID     int64 `db:"user_id"  
    Email  string `db:"email"  
}  
  
// sqlx maps the columns automatically  
users := []User{}  
err := db.Select(&users, "SELECT user_id, email FROM users WHERE active = true")
```





# Connection Pooling for Production

## Managing the Socket Pool

As we learned with HTTP, opening a new database connection for every query is a performance killer. *pgxpool* manages a set of "warm" connections that are reused across goroutines.

### The Strategy:

You must configure your pool limits based on your database's *max\_connections* setting. If your app scales to 50 instances and each tries to open 100 connections, you will crash your Postgres instance.

```
● ● ●  
  
config, _ := pgxpool.ParseConfig(connStr)  
config.MaxConns = 20  
config.MinConns = 5  
  
pool, err := pgxpool.NewWithConfig(context.Background(), config)
```





# The Impedance Mismatch

## Handling Nulls and Custom Types

The "Impedance Mismatch" between SQL and Go is most visible when dealing with *NULL*. A database column can be *NULL*, but a Go *string* or *int* cannot. If you attempt to scan a *NULL* value into a standard Go variable, the driver will return an error, potentially breaking your request flow.

### **The Handling Strategies, depending on your architectural goals:**

1. **Pointers:** Use *\*string* or *\*int* in your struct. If the value is *NULL*, Go sets the pointer to *nil*. This is simple but requires constant nil-checks to avoid panics.
2. **Null Types:** Use *sql.NullString* or the *pgtype* package. These are explicit wrappers (e.g., *Valid: true/false*) that force you to handle the null case safely.
3. **Database Coalescing:** Handle it in the query. Use *COALESCE(column, "")* to ensure the database always returns a zero-value instead of *NULL*, keeping your Go structs clean.





# Security is Non-Negotiable

## Prepared Statements & SQL Injection

Never, ever use `fmt.Sprintf` to build a query with user input. This is how SQL Injection happens.

Always use Parameterized Queries. By using placeholders like `$1` or `?`, you send the query template and the user data separately. The database engine treats the input strictly as data, not as executable code. As a bonus, Postgres can pre-compile these templates, making repeated queries faster.



```
// DANGEROUS: String formatting (SQL Injection Risk)
// If input is: 123; DROP TABLE users;--
// Your DB will execute: SELECT name FROM users WHERE id = 123; DROP TABLE users;--
badQuery := fmt.Sprintf("SELECT name FROM users WHERE id = %s", userInput)
db.Exec(ctx, badQuery)

// SECURE: Parameterized query
// The driver ensures userInput is treated ONLY as a string/ID.
const query = "SELECT name FROM users WHERE id = $1"
row := db.QueryRow(ctx, query, userInput)
```





# Managing Complex Inputs

## Named Queries for Readability

When an *INSERT* has 20 fields, keeping track of \$1, \$2, ... \$20 is a nightmare. *sqlx* allows for Named Queries, where you pass a struct and the driver maps the fields to the placeholders.

The Result: Your code remains readable even as your data models grow. It prevents the "Off-by-one" error where you accidentally swap a user's *email* with their *bio*.

```
● ● ●  
_, err := db.NamedExec(`  
    INSERT INTO users (first_name, last_name, email)  
    VALUES (:first_name, :last_name, :email)`, user)
```





# Summary:

- **Explicit SQL:** Prioritize raw SQL over "Magic" ORMs for clarity.
- **Specialization:** Use *pgx* for Postgres-specific performance gains.
- **Convenience:** Use *sqlx* to map results to structs with zero boilerplate.
- **Security:** Use parameterized queries (\$1, \$2) to prevent injections.

**Question:** Do you prefer the safety of database/sql or the performance features of a driver-specific library like pgx?

