

• JANUARY 2026 SERIES

# FROM GO BUILD TO GO RUN

GOLANG 2026 - NIV RAVE

## #56

# TRANSACTIONS & CONNECTION TUNING

ENSURING DATA INTEGRITY AND PERFORMANCE AT SCALE





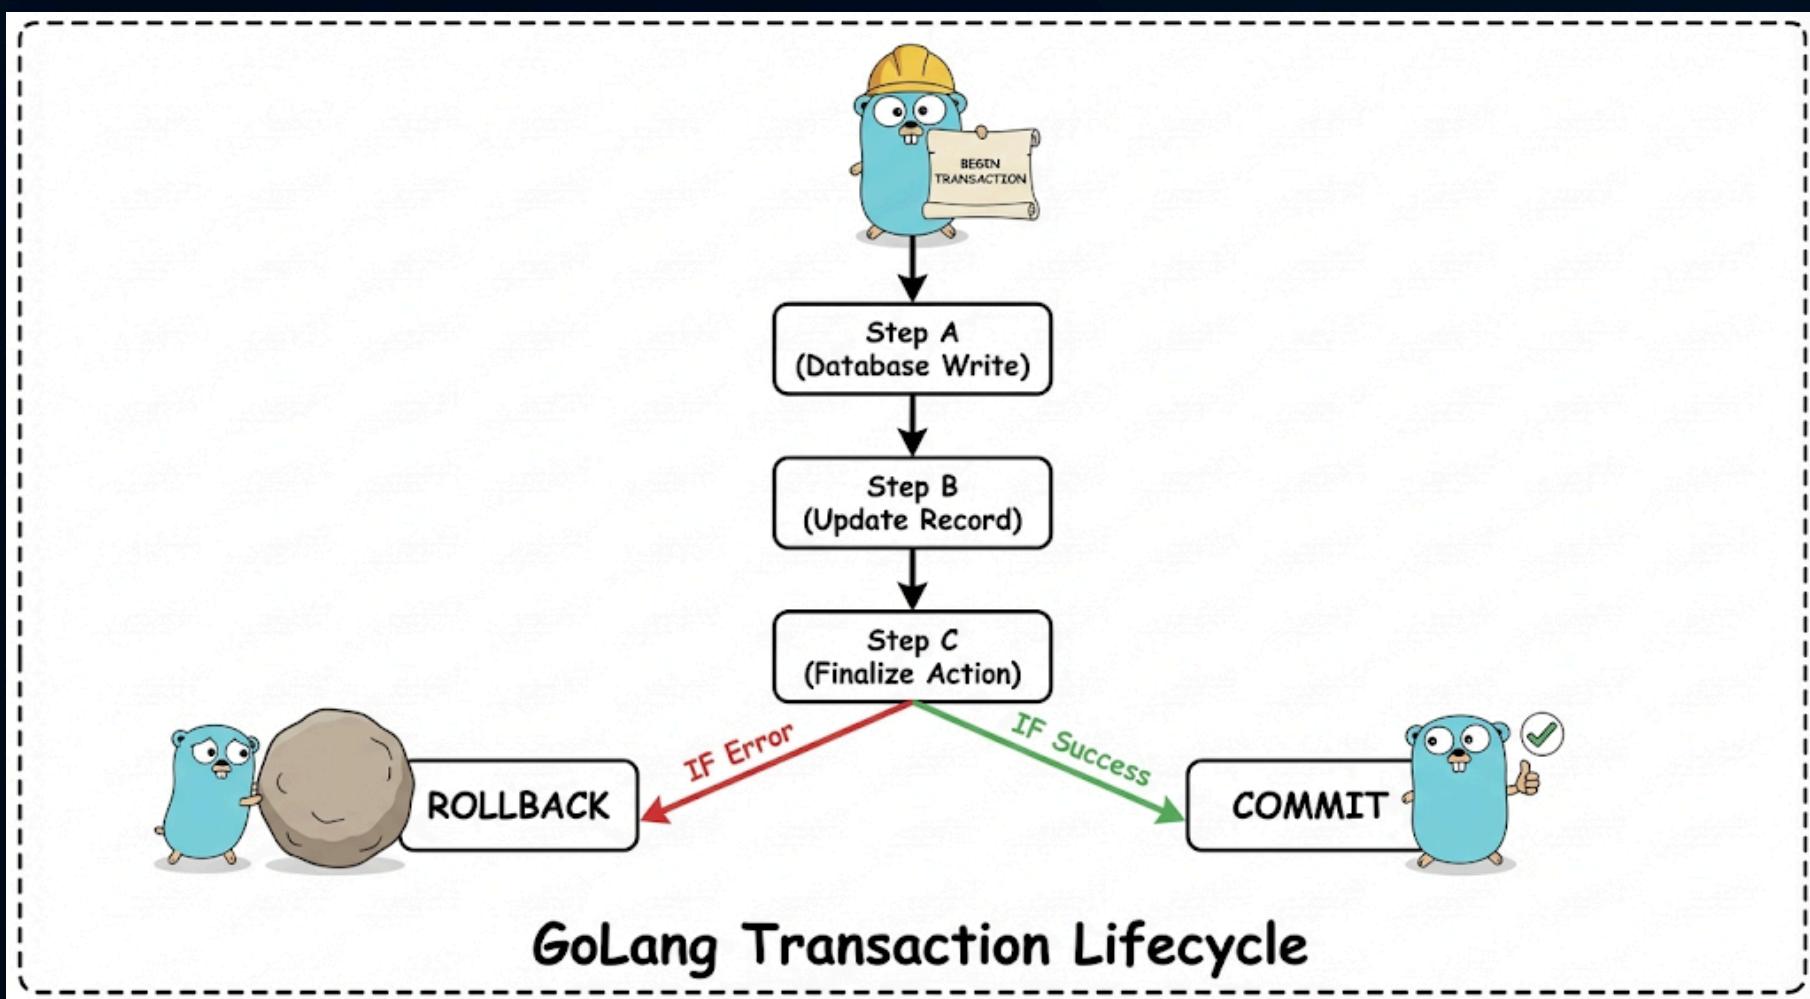
# The ACID Guarantee

## Atomic Operations: All or Nothing

A transaction allows you to group multiple SQL statements into a single unit of work. If any part of the sequence fails, the entire transaction is rolled back, leaving the database in its original, clean state. This is critical for financial records, inventory management, and any logic where partial updates would result in corrupted data.

### The Golden Rule:

Never perform long-running "Business Logic" (like calling an external API) inside a transaction. Keep transactions as short as possible to minimize row locking, which can lead to performance bottlenecks and deadlocks in a high-concurrency system.





# Implementing Transactions in Go

## The Begin-Commit-Rollback Pattern

Go's *database/sql* and *pgx* provide a clear API for transactions. The most important practice is using *defer tx.Rollback()*. If your function returns early due to an error, the rollback is triggered automatically. If you call *tx.Commit()*, the rollback becomes a "no-op," ensuring your data is safe either way.



```
func TransferFunds(ctx context.Context, db *pgxpool.Pool, from, to int, amount float64) error {
    tx, err := db.Begin(ctx)
    if err != nil {
        return err
    }
    // Automatically roll back if we return before committing
    defer tx.Rollback(ctx)

    // 1. Deduct from sender
    if _, err := tx.Exec(ctx, "UPDATE accounts SET bal = bal - $1 WHERE id = $2", amount, from); err != nil {
        return err
    }

    // 2. Add to receiver
    if _, err := tx.Exec(ctx, "UPDATE accounts SET bal = bal + $1 WHERE id = $2", amount, to); err != nil {
        return err
    }

    // Success! Persist the changes
    return tx.Commit()
}
```





# Isolation Levels

## Controlling Visibility

Isolation levels define how and when changes made by one transaction become visible to others. Postgres defaults to Read Committed, which is usually what you want. However, for ultra-sensitive operations like generating a bank statement, you might need Serializable to prevent "Phantom Reads".

### The Trade-off:

Higher isolation levels provide more consistency but decrease performance by increasing the number of locks. As a senior, you must choose the lowest level of isolation that safely maintains your business requirements.

SQL ISOLATION LEVELS: PERFORMANCE vs. CONSISTENCY TRADEOFF

ISOLATION LEVEL	PERFORMANCE (Speed)	CONSISTENCY (Data Integrity)	ANOMALIES PREVENTED
READ UNCOMMITTED	HIGHEST 	LOWEST 	None (Dirty Reads Allowed)
READ COMMITTED	MEDIUM-HIGH 	MEDIUM-LOW 	Dirty Reads
REPEATABLE READ	MEDIUM-LOW 	MEDIUM-HIGH 	Dirty & Non-Repeatable Reads
SERIALIZABLE	LOWEST 	HIGHEST 	All (Dirty, Non-Repeatable, Phantom Reads)





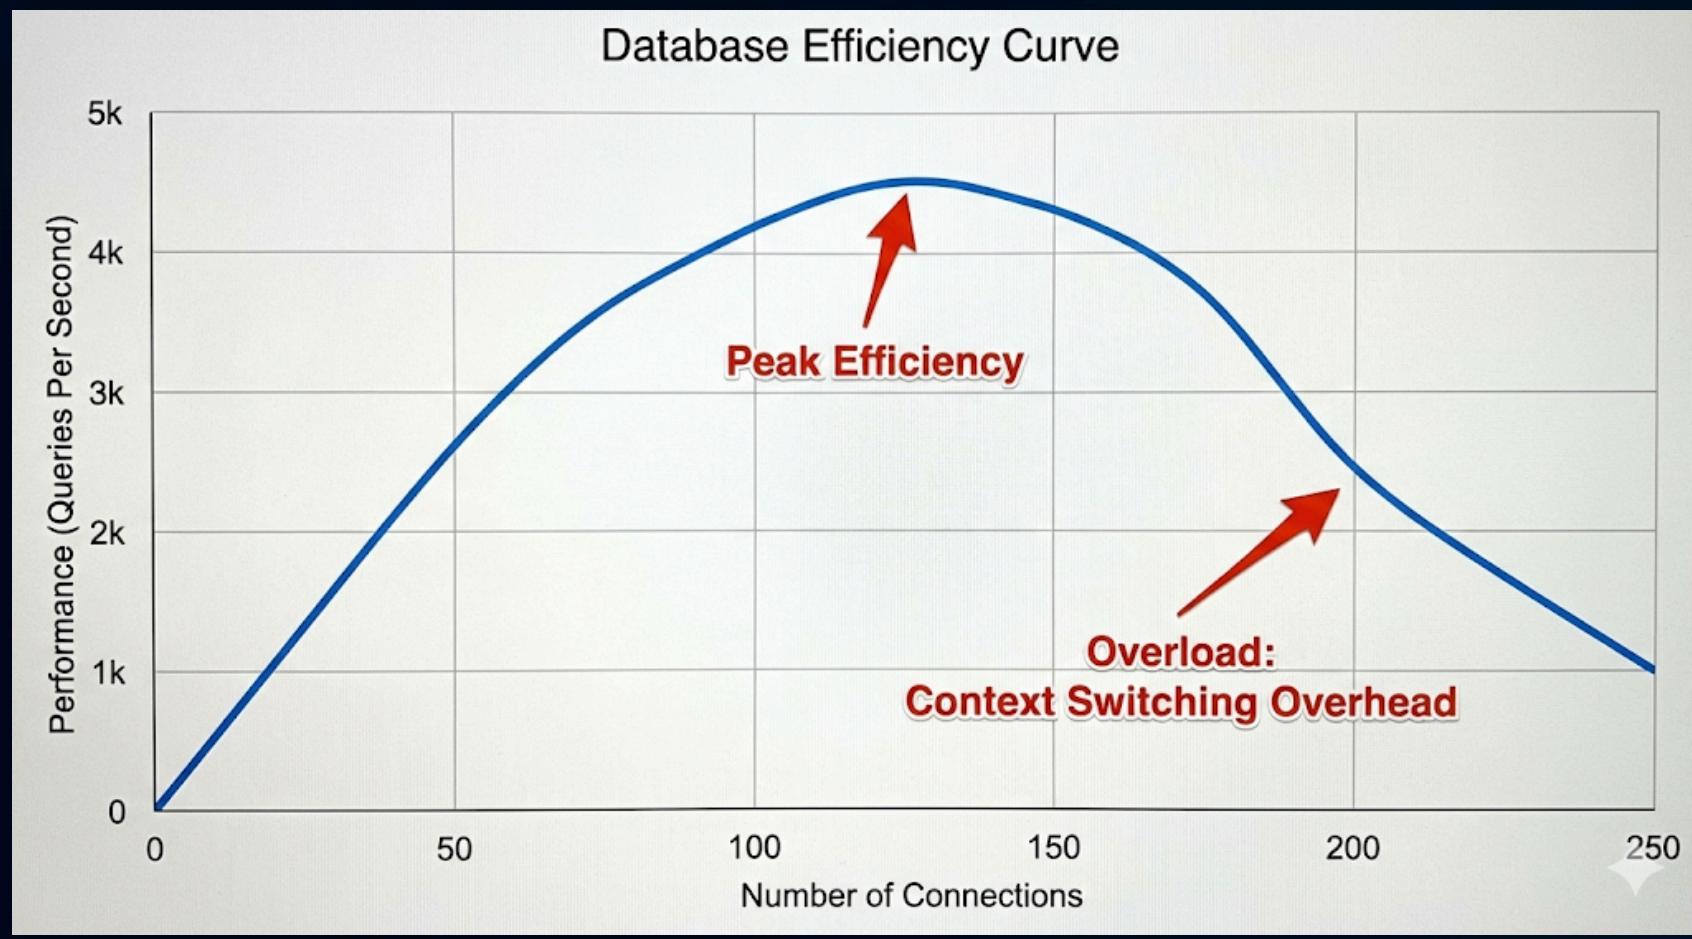
# Connection Pool Tuning

## The "Sweet Spot" of Concurrency

A common mistake is thinking "more connections = more speed". In reality, Postgres uses a process-per-connection model. Too many active connections lead to "Context Switching" overhead, which actually slows down your database.

### The Strategy:

Set your *MaxConns* based on your hardware. A good starting formula is  $(\text{num\_cores} * 2) + \text{effective\_spindle\_count}$ . If you have hundreds of microservices, use a tool like *pgbouncer* to manage connection multiplexing effectively.





# Monitoring Connection Health

## Deadlines and Life Cycles

Connections shouldn't live forever. Network issues can leave "ghost" connections that take up slots in your pool. You need to configure lifetimes and idle timeouts to keep the pool fresh and responsive.

```
● ● ●

config, _ := pgxpool.ParseConfig(os.Getenv("DATABASE_URL"))

// Max time a connection can be reused
config.MaxConnLifetime = 30 * time.Minute

// Max time a connection can sit idle before being closed
config.MaxConnIdleTime = 5 * time.Minute

// Time limit to establish a new connection
config.ConnConfig.ConnectTimeout = 5 * time.Second

pool, _ := pgxpool.NewWithConfig(context.Background(), config)
```





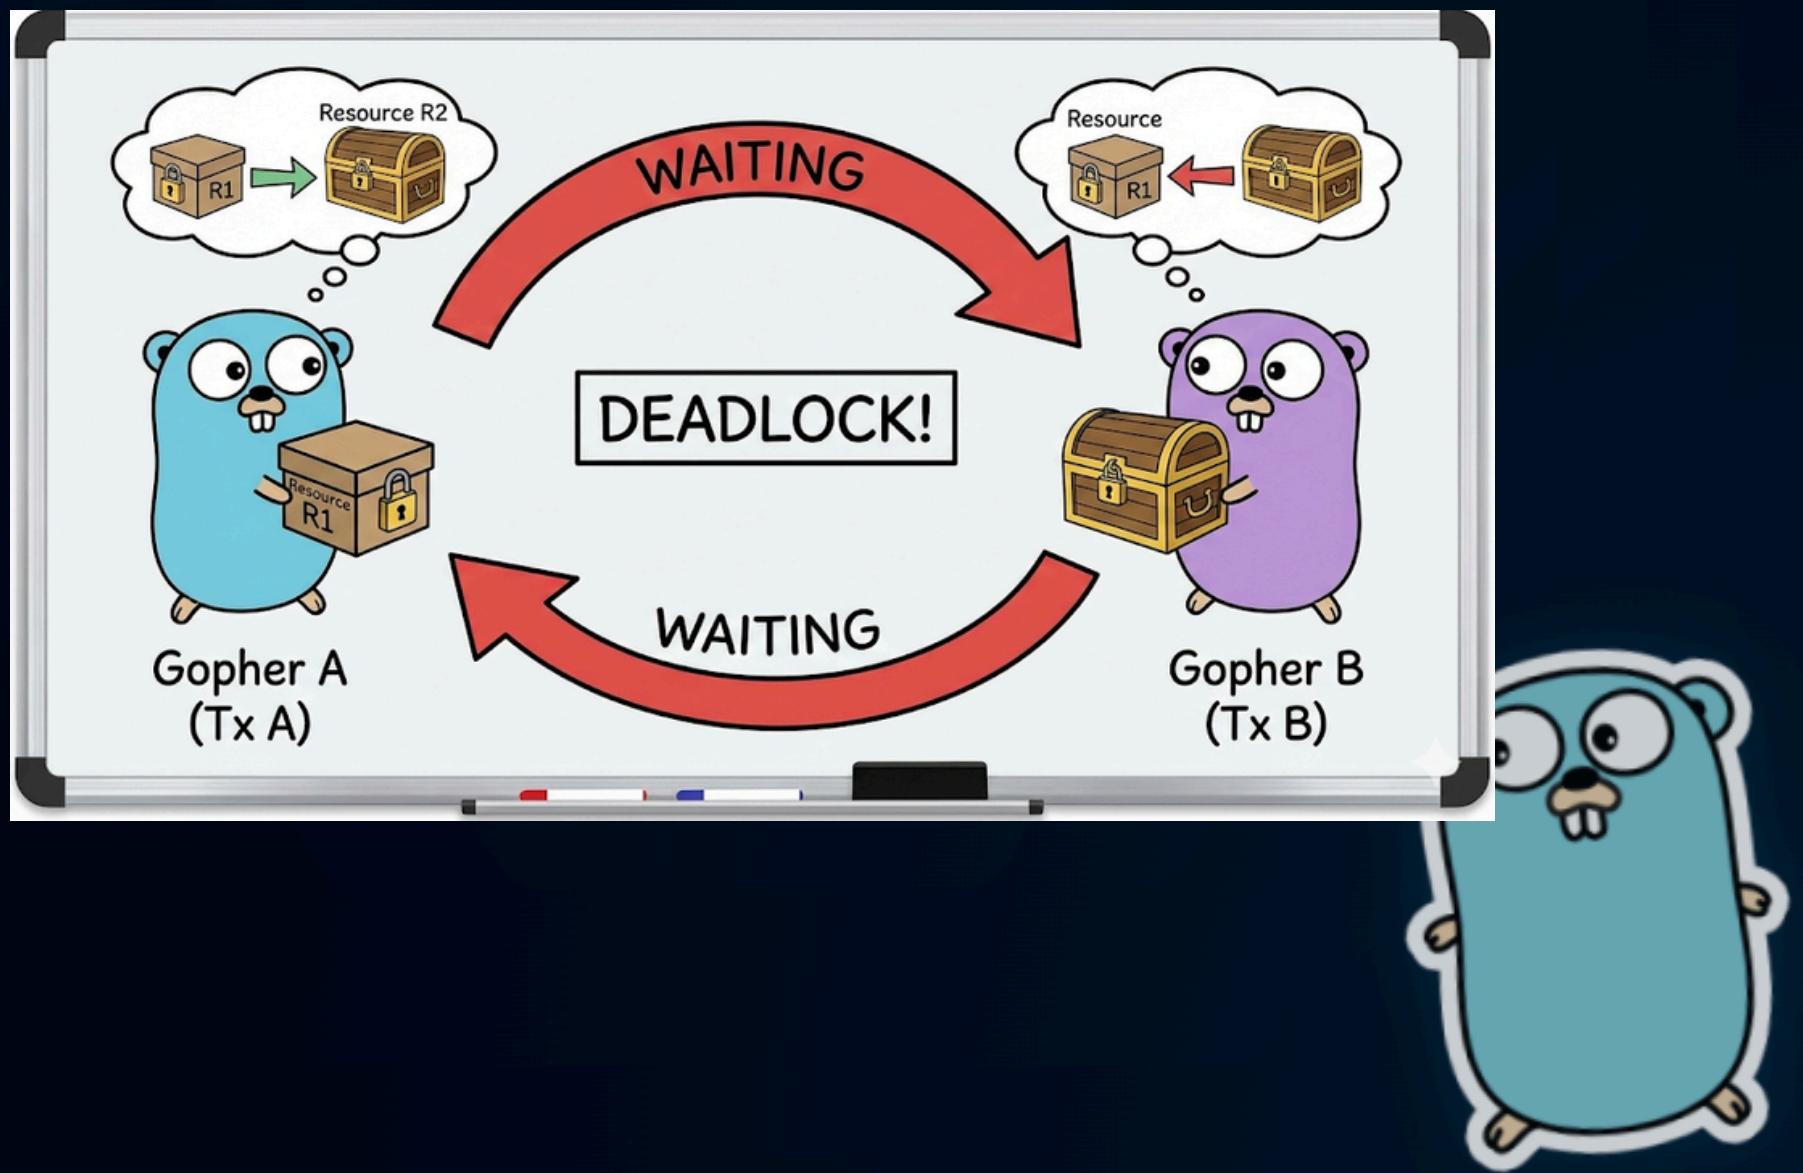
# Dealing with Deadlocks

## Order of Operations Matters

A deadlock occurs when Transaction A holds a lock on Row 1 and wants Row 2, while Transaction B holds Row 2 and wants Row 1. Neither can move. Postgres will eventually detect this and kill one of them, but it's better to avoid it entirely.

### The Prevention:

Always update your tables in the same order across all parts of your application. If every function updates *Users* then *Accounts*, a deadlock is mathematically impossible. This simple convention is the hallmark of a senior-led database architecture.





# Observability for SQL

## Tracking Slow Queries

A database is a black box until you add instrumentation. You should log every query that takes longer than a specific threshold (e.g., 200ms).

This allows you to find missing indexes or unoptimized joins before they cause a production outage.

### The Strategy:

Use *pgx* middleware or interceptors to automatically inject Trace IDs into your SQL comments. When you see a slow query in your Postgres logs, you can instantly find the exact Go function and HTTP request that triggered it.





# Summary:

- **Transactions:** Use `defer tx.Rollback()` to ensure "All or Nothing" safety.
- **Isolation:** Stick to *Read Committed* unless you have a specific reason not to.
- **Tuning:** Leaner connection pools are often faster than bloated ones.
- **Consistency:** Maintain a strict order of operations to prevent deadlocks.

**The Next Step: Tomorrow we enter the world of optimization: Performance Profiling with pprof.**

